

TSN Task Report

Training:

1. RGB model:
 - a. We use a pretrained ResNet18 model, modified to have 25 output classes as in our dataset.
 - b. No. of segments = 4
2. Optical Flow model:
 - a. We also initialize with the pretrained version of ResNet18, because it was shown in the literature that RGB values can correlate with Optical Flow values
 - b. While we are not 100% able to confirm this due to time constraints, we believe that initializing the first layer with the averages learned weights of the RGB model leads to more efficient learning
 - c. No. of segments = 4
3. Notes on training:
 - a. As would be expected, the preinitialization from ImageNet benefits the rgb model much more than the optical flow model, especially at the beginning.

Results:

- We faced problems managing the large memory needs of the different models, especially when training all of them and trying to keep them together in memory for comparison and fusion.
- RGB model reaches 75% accuracy with just one epoch, optical flow (6 epochs) without rgb init 61%, with it 63%. (In the best runs)
- RGB model stays at 75% best accuracy, more epochs lead to overfitting. Loss drops near 0 on the training set, but error gets higher on the test set. We would need more time for testing more parameters and transformations on the data.
- RGB model gets higher results because background information helps a lot in the case of our simple classes

Suggestions for improvement:

- Use more data or apply more transforms and data augmentation to avoid quick overfitting (Training keeps improving on training set, but stagnates on test set)
- Using more segments: We use 4 segments per video, however, in the literature they can use up to 64, so trying more segments would be a good idea
- Trying other baseline models, or deeper versions of ResNet

LOGS from training on one epoch for RGB, 6 for optical Flow in both of its modes:

=====STARTING RGB TRAINING=====

starting epoch for RGB training: 0

/usr/local/lib/python3.7/dist-packages/torchvision/io/video.py:159: UserWarning: The pts_unit 'pts' gives wrong results and will be removed in a follow-up version. Please use pts_unit 'sec'.

"The pts_unit 'pts' gives wrong results and will be removed in a "

Batch 0: loss = 3.4926931858062744

Batch 10: loss = 1.5362038612365723

Batch 20: loss = 0.930392861366272

Batch 30: loss = 0.5449123382568359

Batch 40: loss = 0.32985010743141174

Batch 50: loss = 0.18782034516334534

Batch 60: loss = 0.16709424555301666

Batch 70: loss = 0.19210827350616455

average loss on epoch 0: 0.6672223179539045

evaluating on validation set

current accuracy: 0.7219827771186829

=====STARTING FLOW TRAINING=====

512

starting epoch for optical flow training: 0

Batch 0: loss = 3.298482894897461

Batch 10: loss = 2.4944210052490234

Batch 20: loss = 2.273702383041382

Batch 30: loss = 2.2339725494384766

Batch 40: loss = 1.8463265895843506

Batch 50: loss = 1.7042677402496338
Batch 60: loss = 1.4837234020233154
Batch 70: loss = 1.7284736633300781
average loss on epoch 0: 1.9706972726186116
evaluating on validation set
current accuracy: 0.45743533968925476
starting epoch for optical flow training: 1
Batch 0: loss = 1.058258295059204
Batch 10: loss = 1.0314205884933472
Batch 20: loss = 1.1468636989593506
Batch 30: loss = 0.9844346046447754
Batch 40: loss = 0.7871050834655762
Batch 50: loss = 1.2888381481170654
Batch 60: loss = 1.0843921899795532
Batch 70: loss = 0.9662673473358154
average loss on epoch 1: 1.1353025285402933
evaluating on validation set
current accuracy: 0.5444504022598267
starting epoch for optical flow training: 2
Batch 0: loss = 0.6169149279594421
Batch 10: loss = 0.9653003215789795
Batch 20: loss = 0.6240341663360596
Batch 30: loss = 0.6107838749885559
Batch 40: loss = 0.7504445910453796
Batch 50: loss = 0.6765245795249939
Batch 60: loss = 0.629477322101593
Batch 70: loss = 0.5399748086929321
average loss on epoch 2: 0.797319540977478
evaluating on validation set
current accuracy: 0.595097005367279
starting epoch for optical flow training: 3
Batch 0: loss = 0.6231139898300171
Batch 10: loss = 0.40583735704421997
Batch 20: loss = 0.6060651540756226
Batch 30: loss = 0.5236485600471497
Batch 40: loss = 0.45296353101730347
Batch 50: loss = 0.36324042081832886
Batch 60: loss = 0.6118367314338684
Batch 70: loss = 0.602269172668457
average loss on epoch 3: 0.5489975770314535
evaluating on validation set
current accuracy: 0.6096444129943848
starting epoch for optical flow training: 4
Batch 0: loss = 0.34588488936424255

Batch 10: loss = 0.37583544850349426
Batch 20: loss = 0.3317928910255432
Batch 30: loss = 0.47672486305236816
Batch 40: loss = 0.31999877095222473
Batch 50: loss = 0.40093663334846497
Batch 60: loss = 0.22433197498321533
Batch 70: loss = 0.38348138332366943
average loss on epoch 4: 0.3584156254927317
evaluating on validation set
current accuracy: 0.615840494632721
starting epoch for optical flow training: 5
Batch 0: loss = 0.27572011947631836
Batch 10: loss = 0.2961585819721222
Batch 20: loss = 0.17818039655685425
Batch 30: loss = 0.250670850276947
Batch 40: loss = 0.12327300012111664
Batch 50: loss = 0.3387334644794464
Batch 60: loss = 0.26593053340911865
Batch 70: loss = 0.1790829598903656
average loss on epoch 5: 0.23902002463738123
evaluating on validation set
current accuracy: 0.6371228694915771
=====STARTING FLOW TRAINING with RGB init=====

512

starting epoch for optical flow training: 0
Batch 0: loss = 3.4189157485961914
Batch 10: loss = 2.7579710483551025
Batch 20: loss = 2.1881024837493896
Batch 30: loss = 1.597906470298767
Batch 40: loss = 1.785117745399475
Batch 50: loss = 1.8426909446716309
Batch 60: loss = 1.6441879272460938
Batch 70: loss = 1.714063286781311
average loss on epoch 0: 2.042806951204936
evaluating on validation set
current accuracy: 0.46713361144065857
starting epoch for optical flow training: 1
Batch 0: loss = 1.2274694442749023
Batch 10: loss = 1.2197353839874268
Batch 20: loss = 0.8888126015663147
Batch 30: loss = 1.1054989099502563
Batch 40: loss = 1.3322105407714844
Batch 50: loss = 0.9872933030128479
Batch 60: loss = 0.8750084042549133
Batch 70: loss = 1.1038854122161865
average loss on epoch 1: 1.1333243981997172

```
evaluating on validation set
current accuracy: 0.5463362336158752
starting epoch for optical flow training: 2
  Batch 0: loss = 0.8103703856468201
  Batch 10: loss = 0.9812977910041809
  Batch 20: loss = 0.6072425246238708
  Batch 30: loss = 0.8291133046150208
  Batch 40: loss = 0.8448845148086548
  Batch 50: loss = 0.690477728843689
  Batch 60: loss = 0.7568544745445251
  Batch 70: loss = 0.5800045728683472
average loss on epoch 2: 0.7958146977424622
evaluating on validation set
current accuracy: 0.5522629022598267
starting epoch for optical flow training: 3
  Batch 0: loss = 0.5998767018318176
  Batch 10: loss = 0.48972102999687195
  Batch 20: loss = 0.4167863726615906
  Batch 30: loss = 0.47728854417800903
  Batch 40: loss = 0.33490923047065735
  Batch 50: loss = 0.4067733883857727
  Batch 60: loss = 0.5109941959381104
  Batch 70: loss = 0.6749112606048584
average loss on epoch 3: 0.5360585661729177
evaluating on validation set
current accuracy: 0.5918642282485962
starting epoch for optical flow training: 4
  Batch 0: loss = 0.552790105342865
  Batch 10: loss = 0.3796640634536743
  Batch 20: loss = 0.4437817335128784
  Batch 30: loss = 0.44314926862716675
  Batch 40: loss = 0.7322526574134827
  Batch 50: loss = 0.5619193911552429
  Batch 60: loss = 0.5024798512458801
  Batch 70: loss = 0.2594705820083618
average loss on epoch 4: 0.39306026140848793
evaluating on validation set
current accuracy: 0.6179956793785095
starting epoch for optical flow training: 5
  Batch 0: loss = 0.19502922892570496
  Batch 10: loss = 0.20212356746196747
  Batch 20: loss = 0.2918771803379059
  Batch 30: loss = 0.20447732508182526
  Batch 40: loss = 0.1685357689857483
  Batch 50: loss = 0.1713617891073227
  Batch 60: loss = 0.10694024711847305
  Batch 70: loss = 0.1959858387708664
```

average loss on epoch 5: 0.23295244336128235

evaluating on validation set

current accuracy: 0.6169180870056152

=====TESTING FUSION=====

RuntimeError: CUDA out of memory. Tried to allocate 76.00 MiB (GPU 0;
14.76 GiB total capacity; 13.46 GiB already allocated; 69.75 MiB free;
13.66 GiB reserved in total by PyTorch)